

The Ferdinand Times

VOL. 1 — FRIDAY, SEPTEMBER 18, 2026 — MEXICO CITY — AS OF 12:05 A.M.

YOUR DAY

Unread tech stack meets a long afternoon block

Overnight letters from TLDR, ByteByteGo, Every, and The Pragmatic Engineer; Swim after the morning catchup; Miami and Discovery fill the later hours.



Ferdinand, four of your core tech letters are sitting unread from Thursday: TLDR AI on Claude merging Cowork into regular chat, ByteByteGo on database migrations that cannot take the shop offline, Every on when heavy token use is worth it, and The Pragmatic Engineer's Matt Pocock episode on AI coding skills.

The morning still opens with News / catchup—Hacker Rank, TLDR, and Product Hunt on the calendar note—then Swim from 6:30 to 8:30. Remote also sent a notice that a new payslip is ready to view.

Looking ahead on the same day, Miami runs 4:00 to 6:30 p.m., Discovery follows until 10:00, and Journaling is at 10:00 with the usual reflection prompts.

TODAY

6:00 A.M.	News / catchup Calendar note lists Hacker Rank, TLDR, Product Hunt.	30 MIN
6:30 A.M.	Swim	2 HR
10:15 A.M.	Food	30 MIN
3:00 P.M.	Food	30 MIN
4:00 P.M.	Miami	2.5 HR
6:30 P.M.	Discovery	3.5 HR
10:00 P.M.	Journaling Ask / write / reflection prompts on the event.	1 HR

THIS WEEK

SAT-SUN — Only Sleep blocks returned for the weekend in this pull; no other appointments found.

WED SEP 24 — Metrópolis live score (On the Town only—not on your calendar). Teatro de la Ciudad Esperanza Iris; Festival de Cine Alemán opening.

WEATHER

24–26° / 13–15° HIGH / LOW

Morning looks partly cloudy with haze and a cool start. The afternoon turns warmer and cloudier, with a real chance of short heavy showers (about 5–25 mm in 24 hours) and lightning, per SMN guidance for the Valle de México. Northeast wind around 10–20 km/h, gusts toward 40 km/h.



SAT	Still cloudy	23–25 / 12–14
SUN	Milder, clearer	24 / 14
MON	Similar mild	24 / —
TUE	Similar	24 / —

NEEDS A LOOK

REMOTE

FYI

Payslip ready to view. Notice says the latest payslip is available in the portal. No amounts printed here.

INBOX

TLDR AI

Claude + Cowork merge, ChatGPT sponsored agents, harness tax

FROM TLDR AI · SEP 17, 2026

Claude Cowork and chat are merging into one Claude. Claude Docs and Slides can now produce documents and presentations that users can edit directly, present straight from Claude, or download as a PowerPoint or PDF. The changes will roll out on Pro and Max plans over the next few weeks, with Team and Free plans following soon. Enterprise admins will hear from Anthropic at least 30 days before anything changes for their organizations.

OpenAI announced Sponsored Agents that let users start conversations with business-sponsored agents after clicking an ad in ChatGPT. It also introduced AI-assisted ad creation in ChatGPT Work, new Ads Manager creative tools, and integrations with HubSpot and Shopify.

Google introduced early access to the Model Context Protocol (MCP) for Google Home, enabling AI agents like ChatGPT to control smart home devices. Users must set up a Google Cloud project and provide MCP details to their chosen AI agent for configuration. This update supports all devices in the Google Home ecosystem and is available initially to US subscribers of Google Home Premium Advanced.



EDITOR'S SUMMARY: TLDR goes on to flag research that models are gaming evaluations, a “harness tax” study finding coding-agent wrappers change cost more than success rate, and Translucence’s proposal for embedded evaluators inside frontier labs. Engineering notes cover Agent Substrate on GKE, Meta’s FLAT multimodal tokens, and Ant Group’s finance-tuned open model. Quick hits include Salesforce’s domain model Koa, a new DeepMind Institute, Grok Build memory, and a Mistral-Mozilla browsing partnership.

BYTEBYTEGO

Migrations at Scale: Changing the Application Engine at 30,000 Feet

FROM BYTEBYTEGO

Consider a hypothetical scenario in which an online shopping platform has become extremely successful. It has outgrown its original database. This causes order histories to load very slowly. Maintenance takes much longer than it used to. And engineers spend way too much time managing the workload and the system’s limitations.

The engineering team behind this system wants to replace the database and migrate the data to a new database that promises better performance. But you cannot take a popular online store offline to make the switch. Customers are always placing orders, changing addresses, and requesting refunds. The shop can’t be shut down just like that.

This scenario offers a glimpse of the main challenge inherent in production migrations. How do you change a working system while still running the service that people depend on? At scale, copying existing data during such migrations can take days. Also, many supporting applications may depend on the component being replaced. Every intermediate step needs to work while ordinary business continues.

In this article, we will look at how migrations work at scale and the key strategies that can help make it as efficient as possible.

EDITOR'S SUMMARY: ByteByteGo frames the rest as a playbook for live cutovers—what component is actually being replaced, how to keep dual systems correct while data copy runs for days, and which sequencing tricks reduce risk. The inbox copy stops where the free email/web open ends; the deeper strategy sections sit behind the publication’s paywall.

blog.bytebytego.com/migrations-at-scale-changing-the

INBOX

EVERY

Why You Should Burn More Tokens

BY LAURA ENTIS

Every CEO Dan Shipper consumes more than three times as many tokens as the next-highest user at the company. That gap worries him: Are the rest of us keeping our AI bills low by thinking too small?

A low token spend isn't necessarily a problem, Dan says, "but it should be a smoke signal, especially on the engineering team, to be like, 'Maybe there's more we can be doing here.'"

The goal isn't tokenmaxxing—torching tokens for sport. Instead, Dan wants to give people room to experiment with radical new strategies and ways of working that, should they pan out, could deliver productivity gains that more than justify the token costs.

In a conversation with head of platform Willie Williams, they settled on a tangible starting point: Identify the engineering team's main bottlenecks, see whether AI can remove them, and only then do a cost-benefit analysis.

EDITOR'S SUMMARY: Entis goes on to describe how Every reviews expensive runs after the fact—what it cost, what it bought, what they learned—rather than freezing experimentation up front. A video lead's open-ended 3D-face project spawned an agent swarm that burned enormous usage and produced a weak model; the rebuild capped subagents and added pause-and-review steps. A separate workflow tip: keep a strong planner model, then let it delegate coding to cheaper models when the task is clear.

every.to/why-you-should-burn-more-tokens

THE PRAGMATIC ENGINEER

AI Skills with Matt Pocock

FROM THE PRAGMATIC ENGINEER

In this episode, Matt and I discuss his unconventional path from working as a voice teacher to becoming a developer and going all-in on technical education. He reveals how communication skills helped him break into tech, why he took an unusual three-days-a-week contract at Vercel, and how he built Total TypeScript through workshops, courses, and lots of free tutorials.

We also explore "strategic coding," and how he uses skills like "grill me" and "wayfinder" to plan, delegate, and course-correct with AI agents. Matt explains his "day shift" and "night shift" approach, why splitting context up can keep agents in their "smart zone," and how concepts from classic software engineering books can guide agents to do better.

EDITOR'S SUMMARY: The write-up then lists concrete takeaways: Pocock's "grill-me" skill (agent interviews you relentlessly), borrowing "tracer bullet" language from The Pragmatic Programmer to steer agents, treating the codebase like it must make sense to a colleague with no memory each morning, and shifting more work to cloud agents that keep running after the laptop closes. He is mixed on strict TDD for agents and argues fundamentals matter more as tactical coding gets automated.

newsletter.pragmaticengineer.com/ai-skills-with-matt-pocock

THE TOWN

RECOMMENDED

Slow Horses, Season 6 – Apple TV

Slow Horses is back this week on Apple TV. The new season picks up with the Slough House crew still cleaning up after the last mess, with Gary Oldman's Jackson Lamb steering a team that is better at surviving than impressing HQ. If you like spy stories that stay funny while the mistakes pile up, this is the one to queue.

War Machine – Netflix

War Machine is a fresh Netflix sci-fi action film with Alan Ritchson as an Army Ranger trainee whose field exercise turns into a fight against something that is not in the training manual. It is built for a loud Friday or Saturday night: short setup, then set pieces.

ON THE TOWN


DeVotchKa at Foro Indie Rocks! – tonight

Denver's DeVotchKa plays Foro Indie Rocks! (Sala B) tonight for a Little Miss Sunshine 20th-anniversary set—live versions of the soundtrack that stuck to that road-trip comedy. Doors are listed around 7:00 p.m. in Roma Norte. Folk, rock, and cabaret textures in one bill; 18+ with ID.


Metrópolis with a new live score – Wed Sep 24

The Festival de Cine Alemán opens Sep 24 at Teatro de la Ciudad Esperanza Iris with Fritz Lang's Metropolis and a new live accompaniment by Quique Rangel and Mike Sandoval (2Contrabajos) plus Denise Gutiérrez of Hello Seahorse!. Two contrabasses and voice against the restored silent film, timed to the movie's centenary year.


PUZZLES



6:00 a.m. — News / catchup, as scheduled.



The unread stack brings friends.



Lane two is now a shared inbox.

MINI CROSSWORD

1	2	3	4	5
6				
7				
8				
9				

ACROSS

- 1 Formal film-era address for women (film shelf energy)
6 What SMN issued ahead of today’s Valle showers
7 Mexico City transit nickname, and a nod to living here
8 What a failing test loves to report
9 Afternoon risk in today’s CDMX forecast

DOWN

- 1 Same as 1-Across — old-movie address that still shows up in film talk
2 Weather-service heads-up before the afternoon chubascos
3 How locals shorthand the city’s rail system
4 Red text in a test log
5 What the sky may deliver after lunch today

DAMES · ALERT · METRO · ERROR · STORM

DEAR TIMES

Should I swim before I open the newsletters?

Dear Times, My calendar gives me thirty minutes called News / catchup, then two hours of Swim. Overnight, four dense tech letters arrived unread. If I start reading, I will miss the water. If I swim first, the pile is still there—and maybe larger—when I get back. What is the least-regret move on a storm-risk Friday?

— Lane Two in Mexico City

Dear Lane Two, Swim first. The letters will not expire in two hours, and the pool slot will. Skim subjects after you dry off, then pick one letter to finish—not all four. Save the rest for the Miami or Discovery blocks if those are the flexible ones.

— The Times

JUMBLE

Unscramble each word. Circled answer boxes spell the bonus. Today’s short weather-and-inbox mashup (4 letters).

SSIMW

KETON

MORST

SEDOC

DOWN · WORD · SWIMS · TAKEN · STORM · CODES

WANTED

Part-time newsletter lifeguard

Must blow the whistle when a fifth unread Substack enters the lane. Whistle optional; floaties mandatory. Apply by surviving Friday’s catchup slot.

Quiet room for Journaling at 10:00 p.m.

Must tolerate reflection questions about motivation and tomorrow’s pivot. Soft lamp preferred; no push notifications.

FOR SALE

Slightly damp “News / catchup” half-hour

Used only between 6:00 and 6:30 a.m. Includes one honest attempt to open TLDR before Swim. Buyer collects from the deep end.

One carefully guarded afternoon labeled Miami

Does not include airfare. Transfers as a 4:00–6:30 p.m. focus block only.

NOTICES

Public works: afternoon shower window

Possible citywide. Carry a light shell if Miami or Discovery runs outdoors. This notice is not a calendar appointment.

Department of Discovery

Extended evening session 6:30–10:00 p.m. Citizens are advised that this is a real calendar hold, not a streaming recommendation.