

The Ferdinand Times

VOL. 1 — WEDNESDAY, SEPTEMBER 23, 2026 — MEXICO CITY — AS OF 1:24 P.M.

YOUR DAY

Model-drop mail piles up; evening holds share 10 p.m.

Unread TLDR AI, ByteByteGo, Every, and The Pragmatic Engineer after Sol and Opus launch day; Checkpoint day overlaps Journaling; Builders & Brews is in mail only.



Ferdinand, four of your core tech letters are unread after a busy model-release window: TLDR AI on GPT-6 Sol and Luna next to Claude Opus 5.5, ByteByteGo on customizing and fine-tuning models, Every's Opus 5.5 vibe check, and The Pragmatic Engineer's Maggie Appleton design-engineering episode.

Luma also sent day-of notes for Builders & Brews CDMX this afternoon at Abraxa Lab in Escandón. That meetup is not on your Google Calendar; Miami and Discovery still hold the later afternoon and evening blocks.

Looking ahead tonight, Checkpoint day and Journaling both start at 10:00 p.m. Redpanda's Streamhouse live session is listed for Tuesday, September 29.

TODAY

6:00 A.M.	News / catchup Calendar note lists Hacker Rank, TLDR, Product Hunt.	30 MIN
6:30 A.M.	Swim	2 HR
10:15 A.M.	Food	30 MIN
3:00 P.M.	Food	30 MIN
4:00 P.M.	Miami	2.5 HR
6:30 P.M.	Discovery	3.5 HR
10:00 P.M.	Checkpoint day Same start as Journaling.	2 HR
10:00 P.M.	Journaling Ask / write / reflection prompts on the event.	1 HR

THIS WEEK

THU SEP 24 — Routine blocks return on the weekday pattern. Metrópolis / Techxturas are On the Town only—not on your calendar.

SAT-SUN — Only Sleep blocks returned for the weekend in this pull; no other appointments found.

TUE SEP 29 — Streamhouse (Redpanda) live invite in mail: "AI agents need the present." Not on calendar.

WEATHER

24–26° / 13–15° HIGH / LOW

Morning looks partly cloudy with haze and a cool start. The afternoon turns warmer and cloudier, with a real chance of heavy showers (about 25–50 mm in 24 hours) and lightning, per SMN guidance for the Valle de México. Some local reports tie extra moisture to Pacific Hurricane Polo. East wind around 10–20 km/h, gusts toward 30–40 km/h.



THU	Shower-prone p.m.	26–28 / 13–15
FRI	Mild, lighter rain	25–27 / 13–15
SAT	Cloudier, more showers	23–25 / —
SUN	Similar mild	24–25 / —

NEEDS A LOOK

ANAGRAMA LABS / LUMA **FYI**
Builders & Brews CDMX is today. Day-of and one-hour reminders for the builder meetup at Abraxa Lab in Escandón. Not on your calendar.

REDPANDA **FYI**
Streamhouse live on Tue Sep 29. Invite for a virtual session on AI agents and streaming data.

INBOX

TLDR AI

GPT-6, Opus 5.5, AI leaders at UN

FROM TLDR AI · SEP 23, 2026

OpenAI introduced GPT-6 Sol and Luna as faster, more affordable counterparts to GPT-6 Astra, bringing advances in coding, factuality, computer use, and professional tasks to lower-cost models.

Anthropic introduced Claude Opus 5.5, saying it matched Claude Fable 5.1 on most work while costing less to run than Opus 5. The model also underwent external evaluations and achieved Anthropic's strongest result to date on its automated behavioral audit.

SWE-BENCH PRO V2 releases with 642 tasks from 11 repositories, correcting previous task errors and optimizing evaluation processes. Performance drops for AI models, with OpenAI GPT-5 and Claude Opus 4.1 only scoring around 23% on the public set, highlighting the benchmark's increased challenge and realism compared to SWE-Bench Verified.

EDITOR'S SUMMARY: TLDR goes on to cover Perplexity training its Computer model from real tool-use sessions, Vinod Khosla's trust-and-task-completion framing for personal AI, and a deeper look at what tasks cost on Opus 5.5. Engineering notes include Google's RRSI method for self-improving agents, hardware-agnostic layers in vLLM, and memory-bounded MoE training. Quick hits flag better GPT-6 prompt caching, Meta's Muse debt to OpenClaw, and Altman and Amodei expected at a UN Security Council AI session.



BYTEBYTEGO

How to Customize a Model to Learn New Tricks

FROM BYTEBYTEGO

Modern language models are skilled at many tasks. They can write code, summarize documents, and answer questions across many subjects. However, an application may still need something the model does not provide consistently. Its answers might omit important details, the way it classifies information might confuse similar categories, or its responses might ignore a particular expected writing style.


These things don't mean that the model is not intelligent or doesn't have knowledge. Often, it has the underlying capability but needs additional help applying it according to specific expectations.

That is the purpose of customizing a model. One way is to follow clearer instructions and provide better information. But when those also leave gaps between the expectation and reality of the model, further training needs to be imparted. Techniques such as LoRA and QLoRA make that training more practical by reducing the resources needed to adapt an existing model.

In this article, we are going to look at the various strategies to customize and fine-tune a model.

EDITOR'S SUMMARY: ByteByteGo then walks through prompting and RAG as first tools, then supervised fine-tuning and RLHF when behavior needs to stick. It explains LoRA as frozen base weights plus small adapters, QLoRA as the same idea with a quantized base to save memory, and closes with a practical training checklist: baseline, clean examples, train/validation/test splits, rank and learning-rate choices, and evaluating the combined model—not just the adapter file.

[blog.bytebytego.com / how-to-customize-a-model-to-learn](https://blog.bytebytego.com/how-to-customize-a-model-to-learn)




INBOX

EVERY

Vibe Check: Opus 5.5 Is Pulling Our Codex Converts Back to Claude

BY KATIE PARROTT



Anthropic's new Opus 5.5 promises most of Fable 5.1's performance at roughly 60 percent less per token. After testing it across coding, design, writing, and consulting work, we found a model that makes that claim surprisingly credible—and occasionally loses the assignment while doing extra credit.

Our builders are already switching back. Kieran Klaassen replaced Fable 5.1 with Opus as his daily driver, while former Claude power users Mike Taylor and Tyler Nishida are reconsidering their moves to Codex. The verdict isn't unanimous: Dan Shipper still prefers Codex for knowledge work and writing, and Opus remains an unreliable finisher when time or precision is tight.

Reach for Opus 5.5 when you're building visual products or want a creative collaborator. Keep Fable for problems too large to inspect easily, and keep Astra or GPT-5.6 Sol nearby when the deliverable has a clock. Whatever you choose, set Opus a budget and a stopping point.

EDITOR'S SUMMARY: Parrott's free preview ends there; the fuller scored write-ups sit behind Every's paywall. A companion Dan Shipper note comparing Sol and Opus the same day also arrived unread, with only a short free open available in the inbox copy used here.

[every.to / vibe-check-opus-5-5-is-pulling-our-codex-converts-back-to-claude](#)

THE PRAGMATIC ENGINEER

Design Engineering with Maggie Appleton

FROM THE PRAGMATIC ENGINEER

What can everyone else learn from designers and design engineers? As it turns out, there's plenty, as I discovered when one of the best design engineers in the industry, Maggie Appleton, came onto the Pragmatic Engineer Podcast. She's a staff research engineer at GitHub Next, where she builds prototypes to explore how software engineers might collaborate with AI in new ways.

We got into designers' work and how their design processes are adapting to and changing with AI. We explore why Maggie starts projects with pens and notebooks, what distinguishes design engineers from other designers, and why understanding engineering constraints leads to better collaboration with engineers.

EDITOR'S SUMMARY: Gergely then lists concrete takeaways: Appleton's path from anthropology into tech, learning React by illustrating Egghead tutorials, notebooks as faster than prompt loops for early ideas, agent-built "jigs" with sliders like a personal Figma, and her term "capability gaslighting" for models that impress once and fail the next day. She also argues planning-by-questionnaire with agents overwhelms human decision bandwidth.

[newsletter.pragmaticengineer.com / design-engineering-with-maggie-appleton](https://newsletter.pragmaticengineer.com/design-engineering-with-maggie-appleton)

THE TOWN

RECOMMENDED

Brothers – Apple TV

Brothers premieres today on Apple TV with the first two episodes. Matthew McConaughey and Woody Harrelson play longtime friends whose bond may be closer than they thought. It is a straight comedy launch if you want something lighter after a model-launch inbox day.

Habeas Corpus – Netflix

Habeas Corpus also lands today on Netflix. The Brazilian crime series follows a law professor who leads students to reopen a conviction and finds her brightest pupil tied to her past. A short soundtrack EP by Antonio Pinto and Gabriel Ferreira is out the same day if you want the score first.

ON THE TOWN


Techxturas Sonoras – Thu Sep 24

Teatro Casa de la Paz in Roma Norte hosts the 10th anniversary of Techxturas Sonoras tomorrow evening. Jairo Guerrero leads a live set that turns Mexican literary memory into electronic soundscapes with LIEEVAL, Leticia García of Cabezas de Cera, Marina de Ita Castillo, and Ruth Cabrera. Poetry, visuals, and sound in one room; listed around 8:00–10:00 p.m.


Metrópolis with a new live score – Thu Sep 24

The Festival de Cine Alemán opens tomorrow at Teatro de la Ciudad Esperanza Iris with Fritz Lang's Metropolis and a new live accompaniment by Denise Gutiérrez of Hello Seahorse! and Dos Contrabajos (Quique Rangel and Mike Sandoval). Two contrabasses and voice against the restored silent film; the screening is listed around 7:00 p.m. and runs about two and a half hours.


PUZZLES



Journaling is on the clock at 10.



Checkpoint day claims the same hour.



One pen. Two finish lines.

MINI CROSSWORD

1	2	3	4	5
6				
7				
8				
9				

ACROSS

1 Blank surface before you choose Sol or Opus for the next coding pass

6 When Discovery starts relative to the Miami block

7 Make amends after skipping a day-of builder meetup

8 Voice quality; also the overall vibe of a model review

9 What you do to a bad fine-tune adapter you no longer trust

DOWN

1 Same as 1-Across – empty board before the model pick

2 Same as 6-Across – the later evening hold after Mia-mi

3 Same as 7-Across – apology energy after missing Builders & Brews

4 Same as 8-Across – register of a vibe-check write-up

5 Same as 9-Across – wipe the failed adapter

SLATE · LATER · ATONE · TENOR · ERASE

DEAR TIMES

Do I switch daily drivers again after one vibe-check?

Dear Times, OpenAI and Anthropic dropped new models within hours of each other. Every's team says Opus 5.5 is pulling some Codex converts back, while Dan still prefers Codex for writing. TLDR is full of Sol, Luna, and Opus notes. If I switch tools every launch day, I never finish anything. Do I pick one stack for the rest of the week, or keep testing?

– Tab Limit in Mexico City

Dear Tab Limit, Pick one primary for seven days and park the other behind a single alternate shortcut. Launch-day mail is designed to make every model feel urgent; a short fixed trial tells you more than a round of tab-hopping. Skim the vibe-checks once, then finish one real task on the choice you already opened this morning.

– The Times

JUMBLE

Unscramble each word. Circled answer boxes spell the bonus. What a builder meetup and a model review both need (4 letters).

MSWIS

□○□□□

NARIS

□□○□□

RTSOM

□□□○□

DOELM

□□□○□

WIRE · MODEL · STORM · SWIMS

WANTED

Part-time latency referee

Must decide whether Sol or Opus wins the afternoon without opening a third tab. Whistle optional; timer mandatory. Apply by closing the comparison spreadsheet.

Quiet room for dual 10:00 p.m. checklists

Must host Checkpoint day and Journaling without starting a third ritual. Soft lamp preferred; no push notifications from Luma.

FOR SALE

Slightly overheated “daily driver” title

Previously held by at least two frontier models this week. Includes one honest attempt to stay loyal through a vibe-check. Buyer collects from the downloads folder.

One unused builder-meetup RSVP energy

Generated by day-of email, never promoted onto the calendar. Does not include Metro directions. Transfers as curiosity only.

NOTICES

Public works: afternoon heavy-shower window

Possible citywide. Carry a light shell if Miami or Discovery runs outdoors. This notice is not a calendar appointment.

Department of Discovery

Evening session 6:30–10:00 p.m. remains a real calendar hold. Citizens are advised that streaming premieres are recommendations, not appointments.