


# The Ferdinand Times

VOL. 1 — SATURDAY, OCTOBER 3, 2026 — MEXICO CITY — AS OF 5:45 A.M.

YOUR DAY

## GDG Probabilidad is tonight at 7; still missing from the calendar

Saturday's Bayes/EDA math club is still mail-only. The calendar is blank aside from Sleep. Amadeus Live at Auditorio overlaps 7:30 if you go out. Hetzner's final warning is still open. Thunderstorms are likely today.



Ferdinand, GDG CDMX's Probabilidad & Estadística Científica session is tonight at 7:00 p.m. CST. The Club de Matemáticas para IA reminder has been in mail since Wednesday and still does not appear on Google Calendar. Topics: random variables, Bayes, and EDA with Pandas, Seaborn, and SciPy. Google Calendar today shows only Sleep—no weekday News/Swim/Miami/Discovery skeleton on Saturday.

Friday's Maps Platform + AI codelab is past. Looking ahead, La Casa de Satoshi on Tuesday, October 6, 6:00–9:00 p.m. (approved Luma) still conflicts with Discovery on the calendar and is still off it. Friday, October 9, stacks a pending Grok Bot CDMX × SpaceX-AI morning (10:00 a.m.–1:00 p.m., host approval not in) with an already-approved Reap × AIBM Agentic Buildathon that evening (5:00–9:00 p.m., Sandbox Hub)—both off calendar against the weekday skeleton. Hetzner's Final Payment Warning from September 25 remains open; overnight GitHub was quiet. Mail also closed a Cincel/Paperless Sports World signing flow and flagged a Telmex Infinitum speed bump (850→1000)—utility FYI only. Pack for thunderstorms if you leave for Probabilidad or Amadeus.

WEATHER

~22.5° / ~14.5° HIGH / LOW  
Cool and humid at compile (~15°C, overcast, 95% humidity). Thunderstorms likely through the day (~16 mm, ~94% precip probability). High near the low 20s. Waterproof layer if Probabilidad or Amadeus wins the evening.

|     |                                |             |
|-----|--------------------------------|-------------|
| SUN | Still stormy; cooler high      | ~20 / ~14   |
| MON | Thunderstorm stretch continues | ~22 / ~14   |
| TUE | Wettest day (~27 mm)           | ~22 / ~14.5 |

NEEDS A LOOK

|                                                                                        |        |
|----------------------------------------------------------------------------------------|--------|
| GDG CDMX                                                                               | ACTION |
| Probabilidad & Estadística – tonight 7:00 p.m. Still mail-only; Bayes + EDA math club. |        |
| HETZNER BILLING                                                                        | ACTION |
| Final warning; services still marked blocked. Sep 25 PDF still open.                   |        |
| LUMA / SPACEXAI                                                                        | ACTION |
| Grok Bot CDMX – Oct 9 morning pending. 10:00 a.m.–1:00 p.m.; host approval not in yet. |        |
| LUMA / LA CASA DE SATOSHI                                                              | ACTION |
| AI-risk talk – still approved, still off calendar. Tue Oct 6, 6–9 p.m.; vs Discovery.  |        |
| REAP LATAM / LUMA                                                                      | FYI    |
| Agentic Buildathon spot approved. Fri Oct 9, 5–9 p.m., Sandbox Hub CDMX.               |        |
| TELMEX INFINITUM                                                                       | FYI    |
| Plan speed notice 850→1000. Utility bump; no action required.                          |        |
| CINCEL / PAPERLESS                                                                     | FYI    |
| Sports World doc fully signed overnight. Signing flow completed.                       |        |

TODAY

|           |                                                                                     |          |
|-----------|-------------------------------------------------------------------------------------|----------|
| —         | Weekend blank<br>Sleep only; no weekday skeleton                                    | CALENDAR |
| 7:00 P.M. | GDG: Probabilidad & Estadística Científica<br>Tonight; Bayes + EDA; not on calendar | MAIL     |

THIS WEEK

SAT OCT 3 – GDG CDMX: Probabilidad & Estadística Científica, tonight 7:00 p.m. CST; Bayes + Pandas/SciPy; not on calendar.  
TUE OCT 6 – ¿La IA nos matará a todos? 6:00–9:00 p.m.; La Casa de Satoshi; approved (not on calendar; vs Discovery).  
WED OCT 7 – Checkpoint day 10:00 p.m.–midnight; overlaps Journaling.  
FRI OCT 9 – Grok Bot CDMX con SpaceXAI, 10:00 a.m.–1:00 p.m.; pending approval.  
FRI OCT 9 – Reap x AIBM Agentic Buildathon, 5:00–9:00 p.m.; Sandbox Hub CDMX; registration approved.

TONIGHT

GDG Probabilidad at 7:00 p.m. if you go (mail-only); Amadeus Live at Auditorio 7:30 overlaps; Interestelar Sinfónico at El Cantoral 5–8 p.m. is the earlier film-concert option; thunderstorm odds high – see On the Town.

DUE

Hetzner final warning / services blocked (still open).

## INBOX

## EVERY

# How to Get Better at AI by Asking AI

BY ARIELLE SHIPPER

*Five prompts that move you from one-off chats to orchestrating specialist agents—with Codex as the coach that grades the habits.*

For most of my working life, the first step of any task began with a blank slate: an empty Google Doc, a fresh email, or a new Excel spreadsheet. AI made that process faster, but the way I accomplished it wasn't all that different. I started with a new chat and prompted it over and over until the output matched my unspoken definition of "done."

When I started at Every in May, I was still working that way. Things look very different today. Last week, I needed to verify the cost of purchasing equity for an employee across multiple grants, strike prices, and vesting schedules. Rather than tackle this complex task in a single thread, I split the project among several subagents that coordinated independently to verify information on signed agreements, run the math, double-check it and flag discrepancies, and return a draft of a Slack message that I could approve with a simple "yes."

Asking Codex to assess where I fell on Mike Taylor and Laura Entis's "Eight Levels of AI Adoption" kick-started leveling up. Codex put me at 5.5: solidly at Level 5, with some Level 6 habits. I was using it deeply, but still tended to give one agent a large job rather than deputizing multiple agents with distinct roles.



**EDITOR'S SUMMARY:** Shipper's five moves: (1) ask your main agent to level you from real session history; (2) turn that grade into tactical next-level recommendations on a finished project; (3) demand plain-language how-tos—orchestrator thread, context packets, skeptic/verifier specialists; (4) schedule a weekly graded report (she went 5.5→7.9 in four months); (5) when the model fails tone or accuracy, force a durable Self Improve skill instead of a one-off fix. Pairs cleanly with tonight's TLDR AI clip on decision models that help agents choose when to act or defer. Full piece free on [every.to](https://every.to); Dots / Dot Sebastian now her quick-admin orchestrator.

## TLDR AI

## Decision models land—Clef, Kev, and the classify-then-act stack

FROM TLDR AI

*Cloudflare's open Clef models speak Jev's decision API; Ollama already ships local pulls. Also: Claude-shaped science, OpenAI safety firings, Muse without in-app ads.*

### INTRODUCING CLEF: OUR OPEN-SOURCE DECISION MODELS, AND NEW RL FINE-TUNING PLATFORM


Clef and Clef-flash are fully Jev-API compatible decision models that help agents programmatically gather context, make decisions, and take actions on tasks. The models are fully open-sourced on Hugging Face under an Apache 2.0 license and can be run locally. Decision models make classifications to help agents decide how to act. They return typed answers with probabilities, allowing agents to programmatically gather context, make decisions, and take actions on tasks, or defer to a human when needed.

### CLAUDE-SHAPED SCIENCE

Matthew Schwartz describes using AI, notably Claude, to tackle "Claude-shaped" scientific problems where LLM capabilities like coding and data parsing excel. This led to BootLoops, a toolkit facilitating quantitative science calculations across diverse fields from ecology to population genetics. Despite initial technical correctness, results required domain experts to refine relevance.

### OPENAI CUTS TIES WITH 3 SAFETY RESEARCHERS, WSJ REPORTS

OpenAI dismissed three safety researchers for sharing confidential information with a third-party AI safety group, according to the WSJ. Additionally, OpenAI shelved the GPT-6.1 Astra launch due to safety concerns.



**EDITOR'S SUMMARY:** Friday's TLDR AI flagship is a decision-model race card: Cloudflare Clef/Clef-flash (Jev-compatible, local), Jared Palmer's Kev family (0.8B–27B, same API), Perplexity's pplx-decider-v1-27b, and Amazon's Strands Decider 2B in quick links. Same-day Ollama mail confirms Clef (27B) and Clef Flash (9B) via ollama pull on v0.35.1+. Secondary clips: Microsoft's streaming MAI-Transcribe topping Artificial Analysis, Google's AIM research agent, Meta Muse's ad-free trust thesis, and One Useful Thing on Dots and swarms. Primary newspaper clip is Clef—useful next to Every's orchestration essay.

## INBOX

## TLDR DEVOPS

## Cognitive load, S3 Vectors pre-filter, faster Cloudflare Containers

FROM TLDR DEVOPS

*Agent sandboxes start in under a second; S3 Vectors finally pre-filters metadata before similarity search; AI-named abstractions still tax code review.*

### AMAZON S3 VECTORS NOW SUPPORTS METADATA PRE-FILTERING FOR HIGHER RECALL ON FILTERED SEARCHES

AWS has introduced metadata pre-filtering for Amazon S3 Vectors to improve recall on filtered queries. The feature evaluates metadata filters before running a similarity search and supports up to 100 filter constraints per query. On highly selective filters, pre-filtering returns up to five times more matching vectors than previous methods. It is available immediately at no additional cost and works on existing indexes without requiring data re-ingestion.

### CLOUDFLARE CONTAINERS, REBUILT TO SCALE AGENT SANDBOXES

Cloudflare has rearchitected its Containers service to handle agent workloads, allowing code to select the image and instance type at runtime. The update introduces a new scheduling policy that starts containers six times faster, with median startup times falling from over four seconds to 648 milliseconds in independent benchmarks. A filesystem snapshots feature is also available in public beta to let workspaces save and restore their state.

### REDUCING THE COGNITIVE LOAD OF AI CHANGES

Unfamiliar abstraction names in AI-generated code add mental translation work during review. The same-morning DevOps roundup also flags four levels of agentic software development—most orgs stuck at 1–2—plus HyperFrames and DBX as agent-friendly tooling worth a glance before Monday's deploy queue.



**EDITOR'S SUMMARY:** Same-morning DevOps sibling to Friday's Starlink TLDR (already printed Oct 2)—fuller treatment here. Also: Cloudflare post-quantum Merkle Tree Certificates CA targeting early 2027 Chrome root store; LLM-driven WAF red-teaming that shipped new SSRF detections; Datadog MCP code-execution cutting tokens while raising accuracy; four levels of agentic software development (most orgs stuck at 1–2). HyperFrames and DBX land as agent-friendly tooling. Primary clips for the paper: S3 Vectors pre-filter + Containers startup win; cognitive-load note as the reviewer's punchline.

## THE TOWN

## RECOMMENDED

### Animals – Netflix, Oct 9

Ben Affleck's Netflix thriller lands next Friday, with Kerry Washington, Gillian Anderson, and Steven Yeun. Affleck has said he used AI on the shoot "for lots of things" without tipping which shots. One week out; pairs with the Oct 9 Buildathon / SpaceXAI stack if the evening stays free after the morning pending clears—or if you skip both and keep Discovery.

### Clef on Ollama – local decision models

Cloudflare's Clef and Clef Flash decision models (Jev-API compatible) are pullable on Ollama as of Friday's mail—Clef at 27B and Clef Flash at 9B on v0.35.1+. Useful next to Every's agent-orchestration prompts if you want classify-then-act loops running locally before Monday's coding blocks.

## ON THE TOWN



### Interestelar Sinfónico – Sat Oct 3, 5:00–8:00 p.m., El Cantoral

Hans Zimmer's *Interstellar* score with a 50-plus orchestra while the film runs in 4K—12th-anniversary special at El Cantoral (Xoco). Starts two hours before GDG Probabilidad, so it is an either/or with the math meetup. Indoor, which helps on a thunderstorm Saturday—and it ends right as Probabilidad begins, so stacking is not realistic.

### Lost Acapulco / Namasenda – Sat Oct 3 night

Lost Acapulco at Teatro Metropolitano 8:30 p.m. (indie rock) if you want post-GDG noise; Namasenda (PC Music / AIFF) debuts LIMBO at Foro Niebla 10:00 p.m.–5:00 a.m. for the late slot. Amadeus Live at Auditorio 7:30 overlaps GDG if you try both—pick one lane. Outdoor-adjacent nightlife is optional under tonight's storm odds.

PUZZLES



MINI CROSSWORD

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 |
| 6 |   |   |   |   |
| 7 |   |   |   |   |
| 8 |   |   |   |   |
| 9 |   |   |   |   |

**ACROSS**

1 Tonight's GDG math theme (probability namesake)

6 Live-hall stand-in for tonight's Amadeus venue energy

7 What you do when waitlists and Luma pendings drag on

8 Pass a law – or put Probabilidad on the calendar finally

9 December visitor (wet-season calendar joke)

**DOWN**

1 Same as 1-Across (word-square twin)

2 Same as 6-Across

3 Same as 7-Across

4 Same as 8-Across

5 Same as 9-Across

BAYES · ARENA · YEARN · ENACT · SANTA

**DEAR TIMES**

**Calendar is Sleep-only, GDG wants 7:00, Amadeus wants 7:30—do I pick math, Mozart, or the couch?**

**Dear Times,** Saturday shows only Sleep. Probabilidad is at 7:00 p.m. in mail. Amadeus Live is at 7:30 at Auditorio. Thunderstorms are likely. Interstellar Sinfónico starts at 5:00. Do I put GDG on the calendar and go, take the film-concert lane, or stay in?

— Wet weekend in Mexico City

**Dear Wet weekend,** Put Probabilidad on the calendar this morning so it stops living only as a reminder. Then pick one evening: math club at 7:00, Amadeus at 7:30, or Interstellar at 5:00—do not stack GDG and Amadeus. If storms win, couch is a valid third choice. Before Monday, also place La Casa (Oct 6) and the Oct 9 Buildathon on the calendar; leave SpaceXAI pending until host approval lands. Hetzner still needs a daytime look, not a Saturday-night block.

— The Times

JUMBLE

Unscramble each word. Circled answer boxes spell the bonus. What GDG Probabilidad still needs on Google Calendar (4 letters).

THAM  
○□□□

NIRA  
□○□□

FIEL  
□○□□

ALBL  
□□○□

MATH RAIN FILE BALL · MAIL

**WANTED**

**One honest Saturday 7:00 p.m. decision.**

Applicants must choose Probabilidad, Amadeus, or couch before thunder peaks. Dual RSVPs that assume teleportation between GDG and Auditorio will be returned to sender.

**Calendar entries for events that already have tickets.**

La Casa (Oct 6) and the Agentic Buildathon (Oct 9) remain Luma-approved and Google-absent. Tonight's GDG Probabilidad also accepted—especially before 7:00 p.m.

**FOR SALE**

**Slightly damp umbrella, pre-opened indoors.**

Conditioned by a ~94% thunderstorm forecast and a blank calendar that somehow still has two conflicting nights. Buyer supplies own Bayes notes.

**Spare “ask the model to grade me” energy (conceptual).**

Inspired by Every's five-prompt coaching loop and a Friday TLDR AI stack of decision models. Does not clear a Hetzner warning. Transfers as Inbox reading only.

**NOTICES**

**Department of Pending Approvals**

Grok Bot CDMX con SpaceXAI (Oct 9, 10:00 a.m.–1:00 p.m.) still awaits host OK. Citizens holding Buildathon tickets that evening should not confuse morning pending with evening approved.

**Municipal weather desk**

Cool start ~15°C; high near 22.5°C; thunderstorm odds high all day into midweek (Tue wettest). Interstellar and Amadeus are indoor; outdoor nightlife is optional. This is weather, not a seat assignment.