

# The Ferdinand Times

VOL. 1 — TUESDAY, OCTOBER 6, 2026 — MEXICO CITY — AS OF 5:45 A.M.

YOUR DAY

## La Casa night: “¿La IA nos matará a todos?” starts at 6:00, right on top of Discovery

Both clipper-engine pull requests, #30 and #32, merged overnight. Cursor says your card didn't go through, twice. Rain arrives by mid-morning and turns heavy from about 3:00 p.m.



Ferdinand, Tuesday runs the standard week-day on Google Calendar: News / catchup at 6:00, Swim 6:30–8:30, Food at 10:15 and 3:00, Miami 4:00–6:30, Discovery 6:30–10:00, then Journaling. La Casa de Satoshi's Luma reminder landed last night: ¿La IA nos matará a todos? Riesgo real, poder y propaganda runs 6:00–9:00 p.m. tonight. It is still not on your calendar, so it collides with the last half hour of Miami and most of Discovery. Decide which one wins before you leave for the pool.

GitHub was busy while you slept. PR #32 (single-run worker observability) picked up a Codex review around 11:35 p.m. and merged at 1:03 a.m. CT; PR #30 (Classic V1 plus template selection) merged at 1:09 a.m. A failed-checks email followed for #30 (pr-readiness red, verify green) but the work is already in. Only #29, the Instagram roadmap doc, is still open. Separately, Cursor couldn't process your subscription payment Monday morning and the retry failed again at 4:26 a.m. today. Hetzner's September 25 final warning still sits in Inbox.

TODAY

|                |                                                                                                                             |                 |
|----------------|-----------------------------------------------------------------------------------------------------------------------------|-----------------|
| 6:30–8:30 A.M. | Swim<br>After News / catchup at 6:00; cloudy and ~15°C, dry until about 8:00                                                | CALENDAR        |
| 4:00–6:30 P.M. | Miami<br>Heaviest rain of the day lands right in this block; Food at 10:15 and 3:00                                         | CALENDAR        |
| 6:00–9:00 P.M. | La Casa de Satoshi – or Discovery<br>Luma talk is approved but off calendar; Discovery runs 6:30–10:00, Journaling at 10:00 | LUMA / CALENDAR |

THIS WEEK

**WED OCT 7** – ByteByteGo AI Evals cohort starts. Checkpoint day 10:00 p.m.–midnight; overlaps Journaling.  
**THU OCT 8** – Standard weekday; drizzle likely in the evening.  
**FRI OCT 9** – Grok Bot CDMX con SpaceXAI, 10:00 a.m.–1:00 p.m.; no approval mail yet.  
**FRI OCT 9** – Reap x AIBM Agentic Buildathon, 5:00–9:00 p.m.; Sandbox Hub; approved (not on calendar).

TONIGHT

La Casa's AI-risk talk at 6:00, or Discovery at 6:30. Also tonight: Rodrigo y Gabriela, Geese and a free Carmina Burana—see On the Town.

DUE

Fix the Cursor payment method; put La Casa on the calendar or let it go; Hetzner still open.

WEATHER

~22° / ~15° HIGH / LOW

Cloudy and about 15°C at compile with patchy rain nearby. The swim window stays mostly dry; light rain starts around 8:00 a.m. and lingers through late morning. A brief break near 1:00 p.m., then heavy rain from roughly 3:00 to 6:00 p.m. and steady rain into the night. Bring the umbrella to Miami and to La Casa.

|     |                              |           |
|-----|------------------------------|-----------|
| WED | Patchy rain, brighter midday | ~23 / ~15 |
| THU | Drizzle by evening           | ~23 / ~14 |

NEEDS A LOOK

|                                                                              |         |
|------------------------------------------------------------------------------|---------|
| LUMA / LA CASA DE SATOSHI                                                    | ACTION  |
| Tonight, 6–9 p.m. Approved, still off calendar; vs Discovery.                |         |
| CURSOR BILLING                                                               | ACTION  |
| Subscription payment declined Monday and again at 4:26 a.m. Update the card. |         |
| GITHUB                                                                       | FYI     |
| PRs #30 and #32 merged ~1 a.m. on just-a-clipper-engine. #29 still open.     |         |
| HETZNER BILLING                                                              | ACTION  |
| Final warning; services marked blocked. Sep 25 notice still open.            |         |
| LUMA / SPACEXAI                                                              | WAITING |
| Fri Oct 9, 10 a.m.–1 p.m. Host approval not in yet.                          |         |
| BYTEBYTEGO                                                                   | FYI     |
| AI Evals in Practice cohort starts tomorrow.                                 |         |

INBOX

EVERY

When Trying to Make AI Better Makes It Worse

FROM EVERY · WORKING OVERTIME, BY KATIE PARROTT

A million words of failed drafts taught Katie Parrott that improving her AI setup sometimes means giving it fewer rules. The German word for it: Verschlimmbesserung, “worsen-bettering.”

Her setup had been simple: a folder per column, each with a voice guide, a style guide and drafts, plus an AGENTS.md telling the model what to read. Then she decided to save everything, connect everything and codify everything.

One style guide grew from around 759 words to 3,855, with eight approved openings, six approved endings and two checklists.

“The model couldn’t tell a record from a rule, so it followed all of them.” One essay’s folder ended up holding 127 drafts and more than a million words. A doc review explained why: concrete rules are easier to verify, so they overwhelm the subtler ones while every checklist still passes.

She archived it all into a Historical folder and rebuilt from an empty desktop. The old voice and style guides ran 6,109 words. The new ones run 513.



**EDITOR’S SUMMARY:** Monday’s Every essay is a context-engineering cautionary tale. Saving every correction into a log, cross-linking 385 files and turning repeated notes into checklist items made drafts “crowded and flat.” Her three new principles: the agent asks before writing any file; every folder stands alone; guides stay short and say “often (not always).” Drafts and notes are “task material or evidence, never new instructions.” Plus one rule for herself: no major context decisions under pressure, including the thrill of a new model. Worth a look next to your own agent memory and skills files.

BYTEBYTEGO

The LLM Blindspot: Why Models Forget What’s in the Middle of Your Prompt

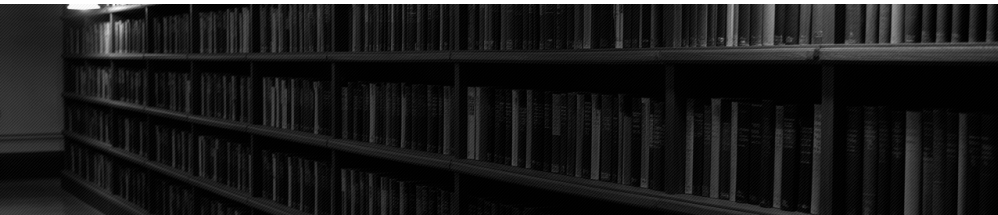
FROM BYTEBYTEGO · MONDAY, OCTOBER 5

The “lost in the middle” effect: a fact can sit inside the context window and still barely shape the answer. Why the beginning and end win, why bigger windows don’t fix it, and what does help.

Imagine a coding assistant given a long set of project documents. One paragraph says audit logs must be retained for 37 days. The assistant writes a cleanup function that deletes them after 30. The rule was there; it was just in the middle.

The 2023 “Lost in the Middle” study found a U-shaped accuracy curve: strongest near the beginning (primacy) and end (recency), weaker in the middle. Causal masking gives early tokens more routes of influence; the end benefits from sitting next to the question.

Larger context windows raise capacity, not uniform reliability. The article separates maximum context size from effective context size for a task.



**EDITOR’S SUMMARY:** Three mitigations. Organize the prompt: state the task and critical constraints up front, label documents with clear boundaries, put the request at the end. Prune context by relevance rather than an arbitrary token budget, keeping exceptions and decisions. Retrieve with RAG, while remembering retrieval can miss passages or reintroduce the same problem if it returns too much. It reads as a companion to the Every essay on the facing column: less, better-placed context beats more.

INBOX

TLDR

iPhone Duo apps open for submission, Google Docs learns Markdown, and the case for ephemeral testing

FROM TLDR · TUESDAY, OCTOBER 6

Apple's two-screen phone gets its App Store lane; Drive and Docs render .md natively; and a short argument for testing a core layer by building throwaway software on top of it.

APPLE INVITES IPHONE DUO-READY APPS

Apple is now accepting submissions of iPhone Duo-optimized apps and games. Xcode 27.1 Release Candidate 1 supports building and testing for the iPhone Duo, App Store pages will show an optimized badge, and recompiling existing apps with Xcode 27.1 adds development support.

MARKDOWN NATIVELY ACROSS DRIVE AND DOCS

Google Docs now supports Markdown files and Google Drive renders Markdown previews. Available to all Workspace customers and personal Google accounts, rolling out gradually over the next two weeks.

EPIHEMERAL TESTING

Instead of guessing what upper layers will need, you simulate them by building them, test the base layer through them, and throw them away when the test is done.



**EDITOR'S SUMMARY:** A practical issue. The Docs/Drive Markdown support matters if your notes, READMEs and agent files live as .md. Also inside: an essay on why plain text is still one of the best technologies we have; Terry Tao on the future of mathematics; Stratechery on a high-severity macOS screen-sharing vulnerability Apple patched last week; tandem solar panels; this year's Physiology Nobel for optogenetics; and a Fitbit Edge preview.

THE TOWN

RECOMMENDED

System Design One: “64 AI Concepts Every Software Engineer Should Know”

Issue #187 is Part 1: generative AI, RAG, AI agents and 14 more concepts, in one glossary-style pass. Pairs well with ByteByteGo’s lost-in-the-middle piece over breakfast after Swim.

Redpanda on evaluating AI agents – today, 10:00 a.m. CT

Redpanda’s CTO and Head of AI talk agent evaluation in a live session at 9:00 a.m. Pacific. A useful warm-up the day before ByteByteGo’s AI Evals cohort begins.

ON THE TOWN



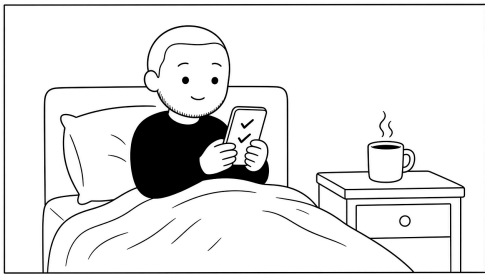
Carmina Después de la Era – Tue Oct 6, 7:00 p.m., Sala Silvestre Revueltas

Centro Cultural Ollin Yoliztli’s new take on Orff’s Carmina Burana, with orchestra and choirs tonight and free entry. The full version with dance plays Wednesday at 7:00 p.m. at Teatro de la Ciudad Esperanza Iris.

Rodrigo y Gabriela and Geese – tonight

The guitar duo plays Polyforum Cultural Siqueiros (doors around 7:00), and Geese bring the Getting Killed Again Tour to Teatro Metropolitán late tonight. Rosana is at Auditorio Nacional. Later: the Interstellar live-score concert at Teatro Xola (Fri) and Beetle-juice, El Musical (Sat, 8:00 p.m.).

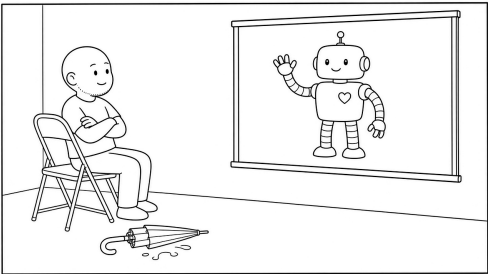
PUZZLES



Two PRs, both merged by 1 a.m.



One talk, no calendar slot.



Will AI kill us all? It waved.

MINI CROSSWORD

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 |
| 6 |   |   |   |   |
| 7 |   |   |   |   |
| 8 |   |   |   |   |
| 9 |   |   |   |   |

- ACROSS
- 1 Emoji for a PR that merged while you slept
- 6 What's left glowing after the campfire
- 7 Misuse, as of an API rate limit
- 8 Tree sap that hardens
- 9 What every daily research email claims to spot
- DOWN
- 1 Same as 1-Across (word-square twin)
- 2 Same as 6-Across
- 3 Same as 7-Across
- 4 Same as 8-Across
- 5 Same as 9-Across

HEART · EMBER · ABUSE · TREND · RESIN ·

DEAR TIMES

A talk about whether AI will kill us all, at the exact hour of my Discovery block —do I go?

Dear Times, Tonight La Casa de Satoshi is hosting a three-hour talk on AI risk, power and propaganda. I'm approved, but it isn't on my calendar and it lands on Discovery. It's also going to pour. Meanwhile Cursor can't charge my card and both my PRs merged at 1 a.m. without me. Where do I put my attention?

– Approved but Undecided in Mexico City

Dear Approved, Swim first, while it's still dry. Then spend two minutes on the Cursor payment method so your tools don't stall mid-week. The PRs are done; celebrate and move on. On tonight, make a decision instead of a collision: if you go, put the talk on the calendar and treat it as tonight's Discovery, since a room full of people arguing about AI risk is discovery. If you stay, archive the reminder. Either way, take the umbrella.

– The Times

JUMBLE

Unscramble each word. Circled answer boxes spell the bonus: what Cursor needs you to update today (4 letters).

RICHA

NIRA

TOBOR

FRATD

CHAIR RAIN ROBOT DRAFT · CARD

WANTED

One calendar slot, tonight 6:00–9:00 p.m.  
¿La IA nos matará a todos? at La Casa de Satoshi is approved in Luma. Discovery has not been consulted. Apply before Miami.

A payment card that goes through.

Cursor subscription seeks working payment method after two declines. Inquire in billing settings.

FOR SALE

Style guide, 3,855 words, eight approved openings.  
Barely followed. Previous owner now uses 513 words and is much happier. Comes with 127 drafts.

Context window, large.

Beginning and end in excellent condition; middle lightly used and sometimes ignored. See ByteByteGo for care instructions.

NOTICES

Department of Merged Work  
just-a-clipper-engine PRs #32 (worker observability) and #30 (Classic V1 + template selection) merged around 1:00 a.m. A late red check on #30 is noted for the record. #29 remains open.

Municipal weather desk

~15°C and cloudy at swim time; light rain from 8:00 a.m.; heavy rain about 3:00–6:00 p.m.; wet through the evening. Wednesday looks brighter at midday.